

# CPSC 340: Machine Learning and Data Mining

Logistic Regression

# Admin

- **Assignment 3:**
  - Due in 9 days
- **Midterm:**
  - March 1<sup>st</sup> in class
  - closed-book, cheat sheet: 1-page double-sided
- **Partners**
  - You can open issues now for hw4, hw5, hw6
- **Office hours**
  - Two extra office hours created for the end of next week, to help with hw3
  - My office hours will move around to accommodate different schedules
  - See calendar for details

# Summary of Last Lecture

## 1. Error functions:

- Squared error is sensitive to outliers.
- Absolute ( $L_1$ ) error and Huber error are more robust to outliers.
- Brittle ( $L_\infty$ ) error is more sensitive to outliers.

## 2. $L_1$ and $L_\infty$ error functions are non-differentiable:

- Finding 'w' minimizing these errors is harder.

## 3. We can approximate these with differentiable functions:




- $L_1$  can be approximated with Huber
  - I was naughty in the demo and didn't do this.
- $L_\infty$  can be approximated with log-sum-exp.

## 4. Gradient descent finds local minimum of differentiable function.

## 5. For convex functions, any local minimum is a global minimum.

# Very Robust Regression

- Consider data with **extreme outliers**:



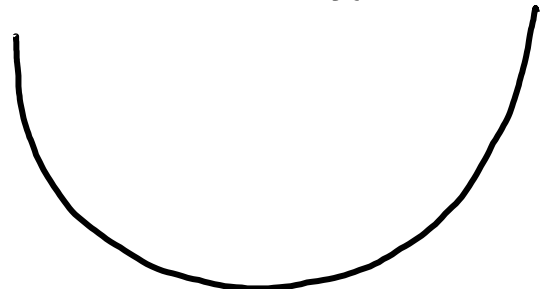
- Non-convex** errors can be **very** robust:
  - Eventually ‘give up’ on trying to make large errors smaller.
- But with non-convex errors, **finding global minimum is hard**.
- Absolute value is the most robust convex error function.**

x x x x x x x x

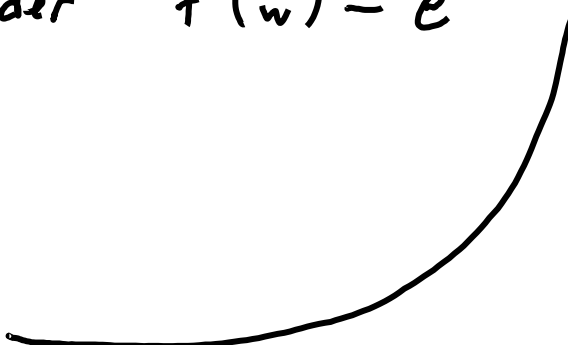
# How do we know if a function is convex?

- Some useful tricks for showing a function is convex:
  - 1-variable, twice-differentiable function is convex iff  $f''(w) \geq 0$  for all 'w'.

Consider  $f(w) = \frac{1}{2}aw^2$  for  $a > 0$ . We have  $f'(w) = aw$   
and  $f''(w) = a > 0$  by assumption



Consider  $f(w) = e^w$



We have  $f'(w) = e^w$   
and  $f''(w) = e^w > 0$   
By definition of exponential function.

# How do we know if a function is convex?

- Some useful tricks for showing a function is convex:
  - 1-variable, twice-differentiable function is convex iff  $f''(w) \geq 0$  for all 'w'.
  - A convex function multiplied by non-negative constant is convex.

We showed that  $f(w) = e^w$  is convex, so  $f(w) = 10e^w$  is convex.

# How do we know if a function is convex?

- Some useful tricks for showing a function is convex:
  - 1-variable, twice-differentiable function is convex iff  $f''(w) \geq 0$  for all 'w'.
  - A convex function multiplied by non-negative constant is convex.
  - Norms and squared norms are convex.

$\|w\|$ ,  $\|w\|^2$ ,  $\|w\|_1$ ,  $\|w\|_\infty$ ,  $\|w\|_1^2$ , and so on are all convex.

# How do we know if a function is convex?

- Some useful tricks for showing a function is convex:
  - 1-variable, twice-differentiable function is convex iff  $f''(w) \geq 0$  for all 'w'.
  - A convex function multiplied by non-negative constant is convex.
  - Norms and squared norms are convex.
  - The sum of convex functions is a convex function.


$$f(x) = \underbrace{10e^w}_{\text{From earlier}} + \underbrace{\frac{\lambda}{2}}_{\text{constant}} \underbrace{\|w\|^2}_{\text{norm squared}} \quad \text{is convex}$$



# How do we know if a function is convex?

- Some useful tricks for showing a function is convex:
  - 1-variable, **twice-differentiable function is convex iff  $f''(w) \geq 0$**  for all 'w'.
  - A convex function **multiplied by non-negative constant** is convex.
  - **Norms** and **squared norms** are convex.
  - The **sum of convex functions** is a convex function.
  - The **max of convex functions** is a convex function.

$$f(w) = \max \{ \underset{\substack{\uparrow \\ \text{convex}}}{1}, \underset{\substack{\uparrow \\ \text{convex}}}{2w}, \underset{\substack{\uparrow \\ \text{convex}}}{w^2} \} \text{ is convex.}$$



# How do we know if a function is convex?

- Some useful tricks for showing a function is convex:
  - 1-variable, twice-differentiable function is convex iff  $f''(w) \geq 0$  for all 'w'.
  - A convex function multiplied by non-negative constant is convex.
  - Norms and squared norms are convex.
  - The sum of convex functions is a convex function.
  - The max of convex functions is a convex function.
  - Composition of a convex function and a linear function is convex.

If 'f' is convex the  $f(Xw - y)$  is convex.

# How do we know if a function is convex?

- Some useful tricks for showing a function is convex:
  - 1-variable, twice-differentiable function is convex iff  $f''(w) \geq 0$  for all 'w'.
  - A convex function multiplied by non-negative constant is convex.
  - Norms and squared norms are convex.
  - The sum of convex functions is a convex function.
  - The max of convex functions is a convex function.
  - Composition of a convex function and a linear function is convex.

- But: not true that composition of convex with convex is convex:

Even if 'f' is convex and 'g' is convex,  $f(g(w))$  might not be convex.

E.g.  $x^2$  is convex and  $-\log(x)$  is convex but  $-\log(x^2)$  is not convex.

# Example: Convexity of Linear Regression

- Consider linear regression objective with error function 'g':

$$f(w) = \sum_{i=1}^n g(w^T x_i - y_i)$$

- Sufficient for 'g' to be convex for 'f' to be convex:
  - Then each term is composition of convex with linear.
  - And sum of convex is convex.

- Examples:

For squared error  $g(r_i) = \frac{1}{2} r_i^2$  so  $g''(r_i) = 1$  and 'f' is convex.

For absolute error  $g(r_i) = |r_i|$  which is a norm so 'f' is convex.

# Example: Convexity of Linear Regression

- Consider linear regression objective with error function 'g':

$$f(w) = \sum_{i=1}^n g(w^T x_i - y_i) + \frac{\lambda}{2} \|w\|^2$$

- Sufficient for 'g' to be convex for 'f' to be convex:
  - Then each term is composition of convex with linear.
  - And sum of convex is convex.
- Same condition applies with  $L_2$ -regularization.

# Classification Using Regression?


- Usual approach to do **classification with regression**:
  - Code  $y_i$  as '-1' for one class and '+1' for the other class.
  - E.g., '+1' means 'important' and '-1' means 'not important'.
- At training time, fit a linear regression model:

$$\begin{aligned} y_i &= w_1 x_{i1} + w_2 x_{i2} + \dots + w_d x_{id} \\ &= w^T x_i \end{aligned}$$


- To predict, we **take the sign**:

$$y_i = \text{sign}(w^T x_i) \begin{cases} \text{Set } y_i = +1 \text{ if } w^T x_i > 0 \\ \text{Set } y_i = -1 \text{ if } w^T x_i < 0 \end{cases}$$

# Classification using Regression




# Classification using Regression





# Classification Using Regression

- Can use regression tricks (basis, regularization) for classification.
- But, usual error functions do weird things:

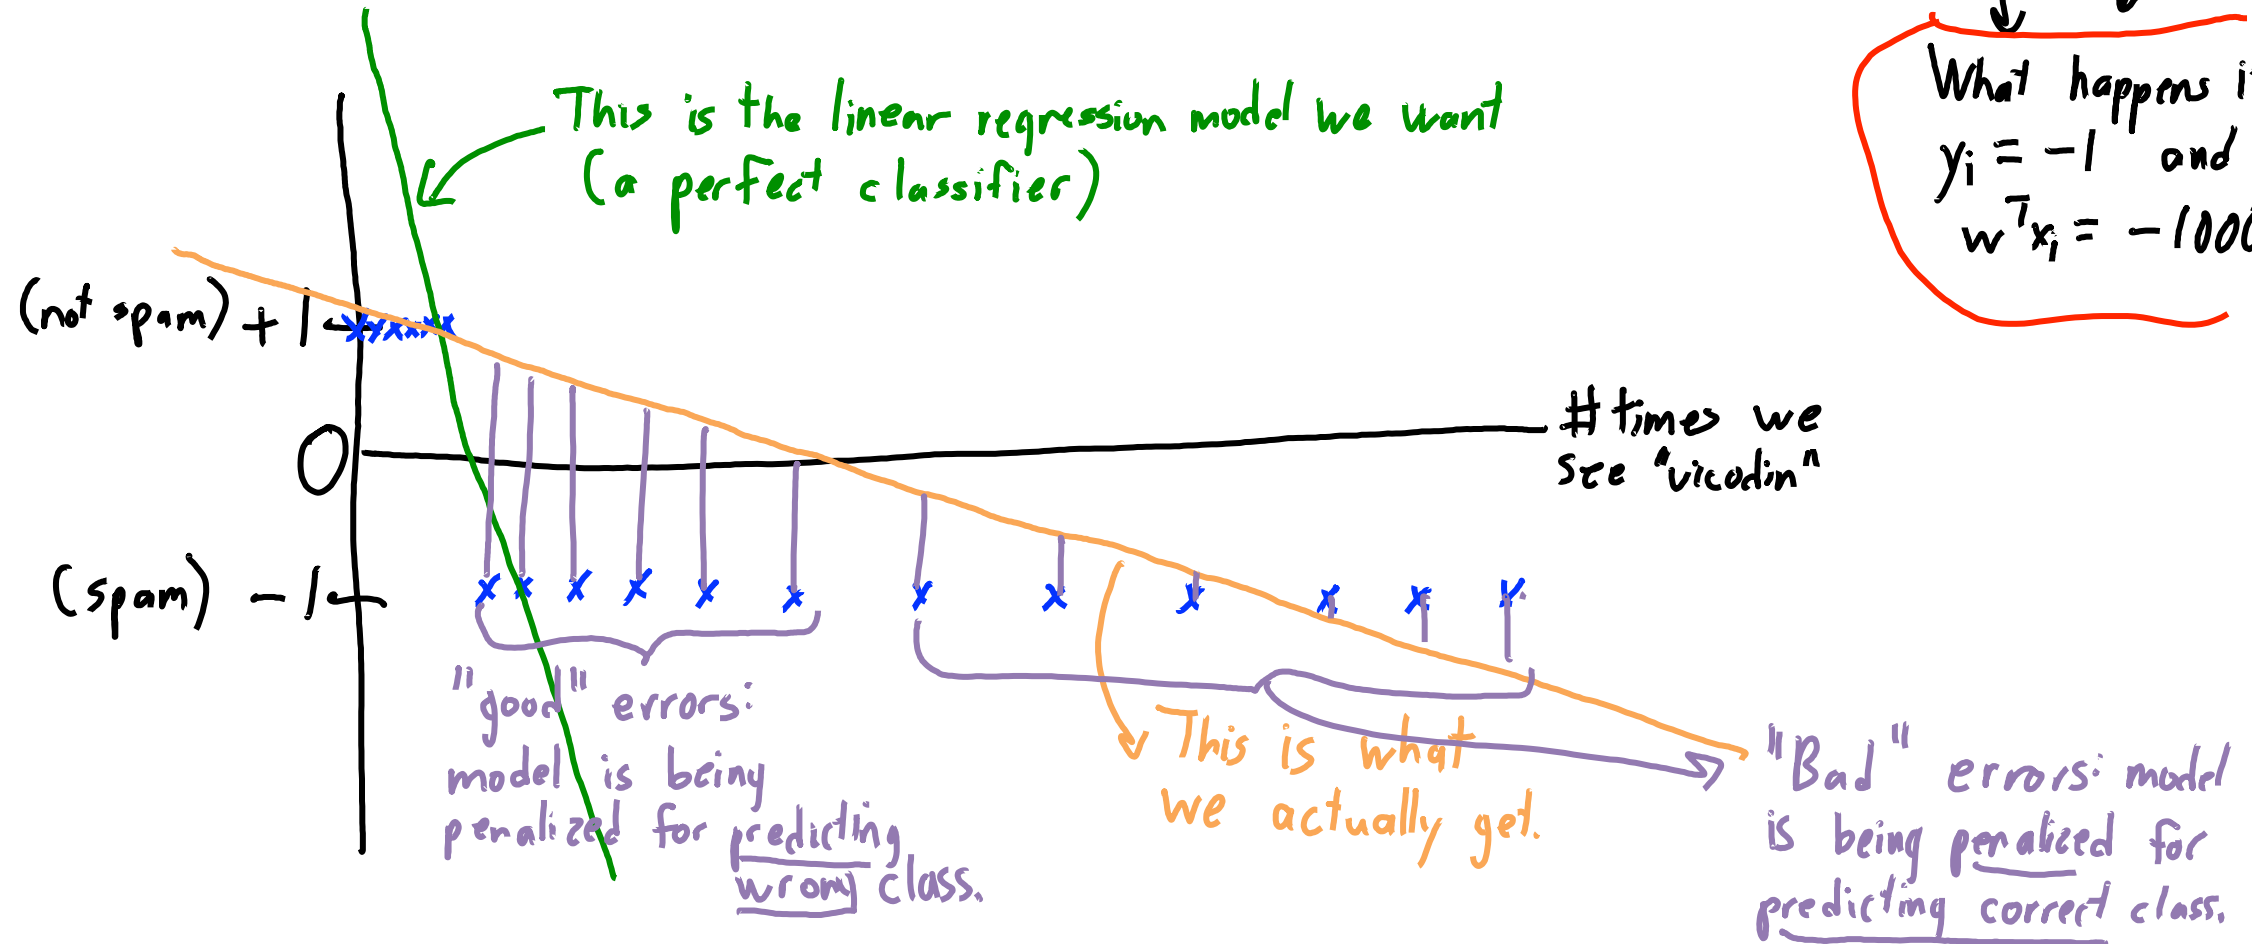


# Classification Using Regression

- What went wrong?
  - “Good” errors vs. “bad” errors.

$$f(w) = \sum_{i=1}^n (w^T x_i - y_i)^2$$

What happens if  
 $y_i = -1$  and  
 $w^T x_i = -1000$ ?




# Classification Using Regression


- What went wrong?
  - “Good” errors vs. “bad” errors.

$$f(w) = \sum_{i=1}^n (w^T x_i - y_i)^2$$


What happens if  
 $y_i = -1$  and  
 $w^T x_i = -1000$ ?




# Comparing Loss Functions



# Comparing Loss Functions




# Comparing Loss Functions



# 0-1 Loss Function


- The 0-1 loss function is the number of classification errors:
  - Unlike regression, in classification it's reasonable that  $\text{sign}(w^T x_i) = y_i$ .
- Unfortunately the 0-1 loss is non-convex in 'w'.
  - It's easy to minimize if a perfect classifier exists.
  - Otherwise, finding the 'w' minimizing 0-1 loss is a hard problem.
  - It's not differentiable, so you don't know "which way to go" in w-space.
- Convex approximations to 0-1 loss:
  - Hinge loss (non-smooth) and logistic loss (smooth).

# Convex Approximations to 0-1 Loss





# Convex Approximations to 0-1 Loss



# Hinge Loss and Support Vector Machines

- Hinge loss is given by:

$$f(w) = \sum_{i=1}^n \max\{0, 1 - y_i w^T x_i\}$$

- Convex upper bound on number of classification errors.
- Solution will be a perfect classifier, if one exists.
- Support vector machine (SVM) is hinge loss with L2-regularization.

$$f(w) = \sum_{i=1}^n \max\{0, 1 - y_i w^T x_i\} + \frac{1}{2} \|w\|^2$$

- Next time we'll see that it “maximizes the margin”.
- Note: it's important that we define  $y$  in  $\{-1, +1\}$  rather than  $\{0, 1\}$ 
  - This allows convenient/compact notation for the loss, as above

# Logistic Regression

- Logistic regression minimizes logistic loss:

$$f(w) = \sum_{i=1}^n \log(1 + \exp(-y_i w^T x_i))$$

- You can/should also add regularization:

$$f(w) = \sum_{i=1}^n \log(1 + \exp(-y_i w^T x_i)) + \frac{\lambda}{2} \|w\|^2$$


- Convex and differentiable: minimize this with gradient descent.

# Probabilistic interpretation

- One can arrive at logistic regression from a completely different viewpoint
- It's in a class called Generalized Linear Models (GLMs)
- You can interpret the logistic function as turning  $w^T x$  into a **probability**:

$$\Pr(y_i = +1) = \frac{1}{1 + \exp(-w^T x_i)}$$

- This function maps the real line to  $[0,1]$  and is “symmetric”
- We set ‘w’ to the maximum likelihood estimate given the data
- Note: don't confuse the logistic loss with the logistic function (below)



# Logistic Regression and SVMs


- Logistic regression and SVMs are **used EVERYWHERE!**
- Why?
  - Training and testing are both fast.
  - It is easy to understand what the weights ' $w_j$ ' mean.
  - With high-dimensional features and regularization, often good test error.
  - Otherwise, often good test error with RBF basis and regularization.
  - Smoother predictions than random forests.

# Linear Models with Binary Features

- What is the effect of a binary feature on linear regression?

Year	Gender
1975	1
1975	0
1980	1
1980	0

Height
1.85
2.25
1.95
2.30



- Adding a bias  $w_0$ , our linear model is:

$$\text{height} = w_0 + w_1 * \text{year} + w_2 * \text{gender}$$

- The 'gender' variable causes a **change in y-intercept**:


If  $\text{gender} = 0$  then  $\text{height} = w_0 + w_1 * \text{year}$

If  $\text{gender} = 1$  then  $\text{height} = w_0 + w_1 * \text{year} + w_2$

new y-intercept

# Linear Models with Binary Features

- What if different genders have different slopes?
  - You can use gender-specific feature (as if  $d=4$ ).
  - But now the two models are completely separate.



Year	Gender
1975	1
1975	0
1980	1
1980	0



Bias (gender = 1)	Year (gender = 1)	Bias (gender = 0)	Year (gender = 0)
1	1975	0	0
0	0	1	1975
1	1980	0	0
0	0	1	1980

$$\text{distance} = w_0 + w_1 * \text{year} \quad (\text{if gender} = 1)$$

$$\text{distance} = w_3 + w_4 * \text{year} \quad (\text{if gender} = 0)$$


separate  
bias

separate slope

# Linear Models with Binary Features

- That trick fits separate ‘local’ variable for each gender.
- To share information across genders, include a ‘global’ version.

Year	Gender
1975	1
1975	0
1980	1
1980	0



Year	Year (if gender = 1)	Year (if gender = 0)
1975	1975	0
1975	0	1975
1980	1980	0
1980	0	1980

- ‘Global’ year feature: influence of time on both genders.
  - E.g., improvements in technique.
- ‘Local’ year feature: gender-specific deviation from global trend.
  - E.g., different effects of performance-enhancing drugs.

*Handwritten notes:* "global" across genders, "local" to gender


$$y_i = w_0 + w_1 * \text{year} + w_3 * \text{year}$$




# Linear Models with Categorical Features

$X =$

Feature 1	Feature 2
0.5	X
3	O
5	O
2.5	$\Delta$
1.5	X
3	$\Delta$
...	...




# Linear Models with Categorical Features



# Linear Models with Categorical Features

$X =$

Feature 1	Feature 2
0.5	X
3	O
5	O
2.5	$\Delta$
1.5	X
3	$\Delta$
...	...



Model 1: only bias  
 $y_i = w_0$

Model 2: bias + feature 1  
 $y_i = w_0 + w_1 x_{i1}$

# Linear Models with Categorical Features

$X =$

Feature 1	Feature 2
0.5	X
3	O
5	O
2.5	$\Delta$
1.5	X
3	$\Delta$
...	...

$$w_0 + w_1 x_{i1}$$

$$w_0 + w_1 x_{i1}$$

$$w_0$$

Model 1: only bias

$$y_i = w_0$$

Model 2: bias + feature 1


$$y_i = w_0 + w_1 x_{i1}$$

Model 3: "local" bias + feature 1

$$y_i = w_{\text{shape}} + w_1 x_{i1}$$

↑ shape

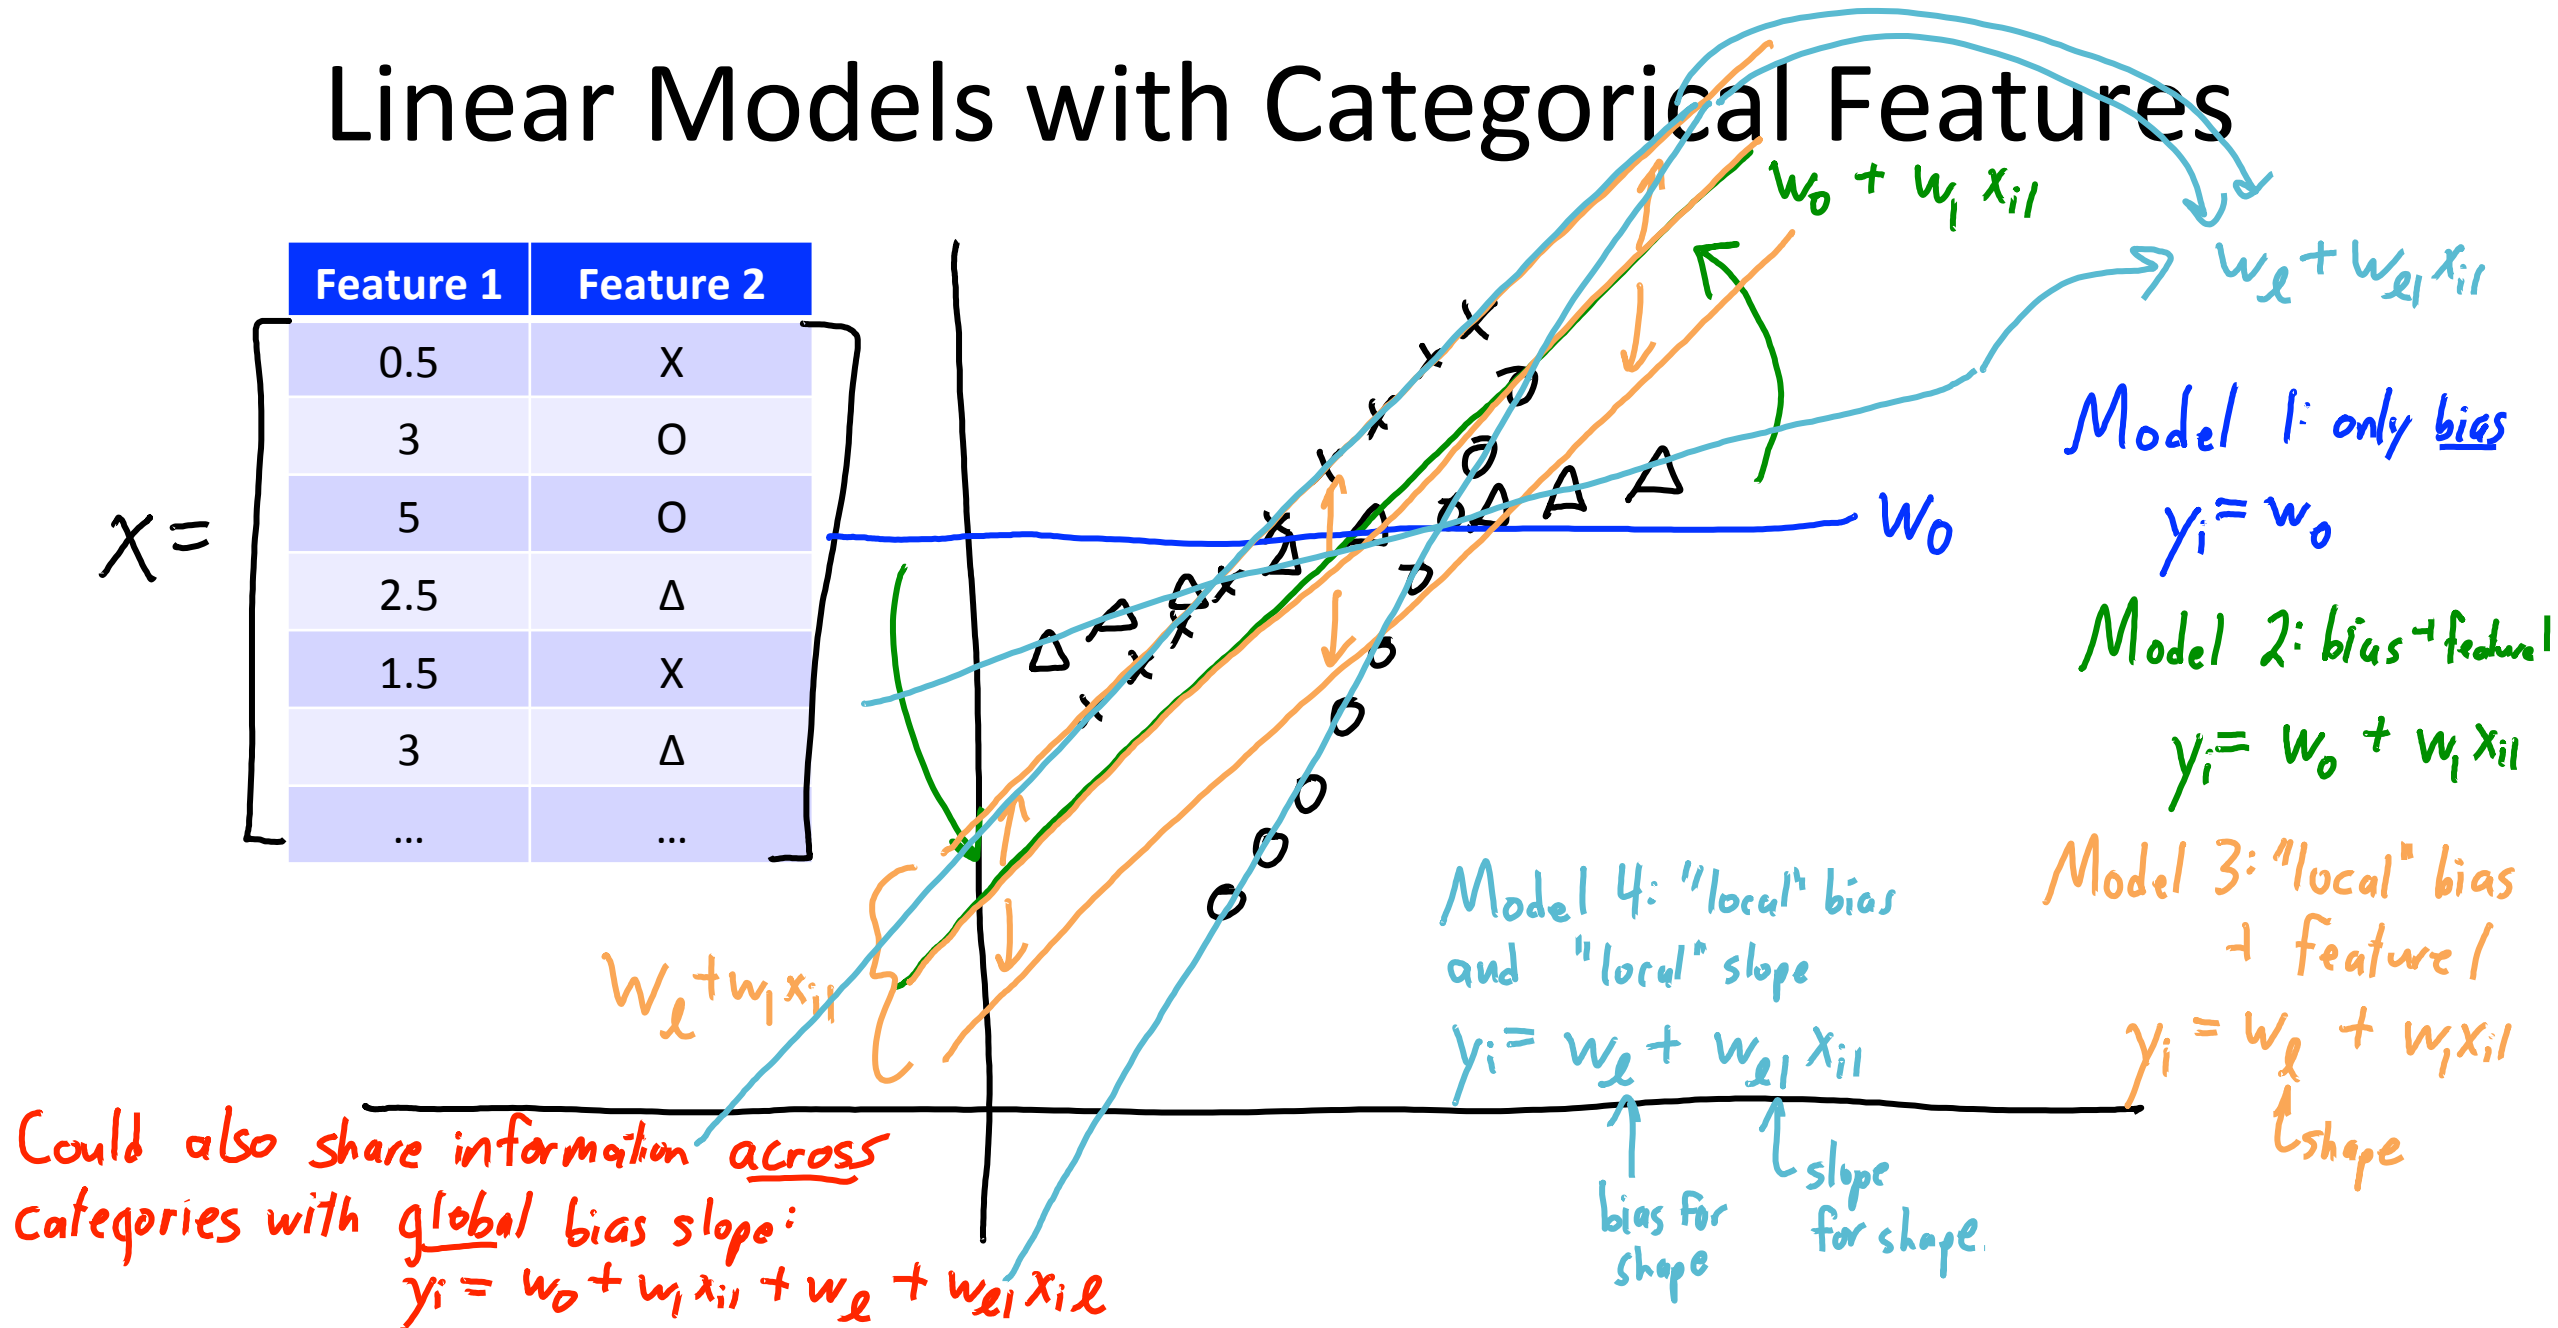
# Linear Models with Categorical Features



# Linear Models with Categorical Features

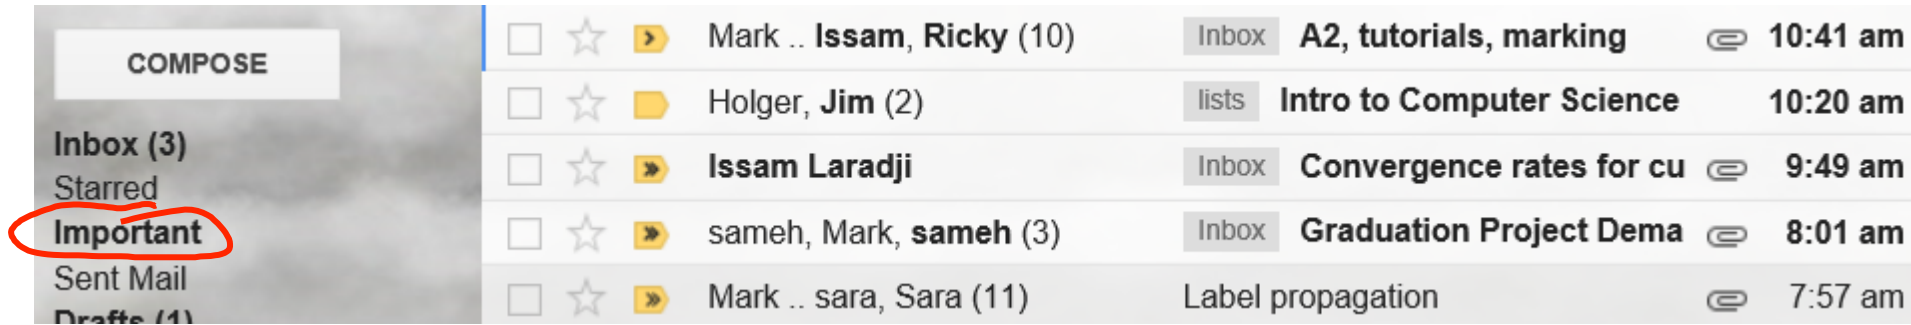
$X =$

Feature 1	Feature 2
0.5	X
3	O
5	O
2.5	$\Delta$
1.5	X
3	$\Delta$
...	...




# Motivation: Identifying Important E-mails

- How can we automatically identify ‘important’ e-mails?



- We have a big collection of e-mails:
  - Mark as ‘important’ if user takes some action based on them.
- There might be some “universally” important messages:
  - “This is your mother, something terrible happened, give me a call ASAP.”
- But **your “important” message may be unimportant to others.**
  - Similar for spam: “spam” for one user could be “not spam” for another.

# The Big Global/Local Feature Table





# Predicting Importance of E-mail For New User

- Consider a new user:
  - Start out with no information about them.
  - Use **global** features to predict what is important to generic user.

$$y_i = \text{sign}(w_g^T x_g)$$

features/weights shared across users

- With more data, update **global** features and **user's local** features:
  - **Local** features make prediction *personalized*.

$$y_i = \text{sign}(w_g^T x_g + w_u^T x_u)$$

features/weights specific to user.

- What is important to *this* user?
- Gmail's system: classification with **logistic regression**.

# Summary

- Convex functions can be identified using a few simple rules.
- Classification using regression works if done right.
- 0-1 loss is the ideal loss, but is non-smooth and non-convex.
- Logistic regression uses a convex and smooth approximation to 0-1.
- Global vs. local features allows 'personalized' predictions.
- Next time:
  - Support Vector Machines

# Bonus Slide: Perceptron Algorithm

- One of the first “learning” is the perceptron algorithm.
  - Searches for a ‘ $w$ ’ such that  $w^T x_i > 0$  when  $y_i = +1$ ,  $w^T x_i < 0$  for  $y_i = -1$ .
- Perceptron Algorithm:
  - Start with  $w^0 = 0$ .
  - Go through examples in any order until you make a mistake predicting  $y^i$ .
    - Set  $w^{t+1} = w^t + y_i x_i$ .
  - Keep going through examples until you make no errors on training data.
- If a perfect classifier exists, this algorithm converges to one.
  - In fact, “perceptron mistake bound” result says that number of mistakes is finite.